


Compiler Design

Lecture 8: Runtime Environments

Dr. Momtazi
momtazi@aut.ac.ir

based on the slides of the course book

Intermediate Representation Generation.



■ IR Generation

- Goal: Translate the program into the format expected by the compiler back-end.

Outline

- **Introduction**
- Runtime Environments
- Encoding in Runtime Environments
- Control-flow

What is IR Generation?

- Intermediate Representation Generation.
- The final phase of the compiler front-end.
- Translate the program into the format expected by the compiler back-end.
- Generated code need not be optimized
 - It will be handled by later passes.
- Generated code need not be in assembly
 - It can be handled by later passes.

Why Do IR Generation?

- Simplify certain optimizations.
 - Machine code has many constraints that inhibit optimization.
 - Working with an intermediate language makes optimizations easier and clearer.
- Have many front-ends into a single back-end.
 - gcc can handle C, C++, Java, Fortran, Ada, and many other languages.
 - Each front-end translates source to the GENERIC language.
- Have many back-ends from a single front-end.
 - Do most optimization on intermediate representation before emitting code targeted at a single machine.

Designing a Good IR

- IRs are like type systems – they're extremely hard to get right.
- Need to balance needs of high-level source language and low-level target language.
- Too high level: can't optimize certain implementation details.
- Too low level: can't use high-level knowledge to perform aggressive optimizations.
- Often have multiple IRs in a single compiler.

Steps

- Runtime Environments
- Three-Address Code IR

Outline

- Introduction
- **Runtime Environments**
- Encoding in Runtime Environments
- Control-flow

An Important Duality

- Programming languages contain high-level structures
- The physical computer only operates in terms of several primitive operations

An Important Duality

- High-level structures in programming languages:
 - Functions
 - Objects
 - Exceptions
 - Dynamic typing
 - ...

An Important Duality

- Primitive operations of physical computer:
 - Arithmetic
 - Data movement
 - Control jumps

Runtime Environments

- We need to come up with a representation of these high-level structures using the low-level structures of the machine.
- A runtime environment is a set of data structures maintained at runtime to implement these high-level structures.
 - e.g. the stack, the heap, static area, virtual function tables, etc.

Runtime Environments

- Strongly depends on the features of both the source and target language. (e.g compiler vs. cross-compiler)
- Our IR generator will depend on how we set up our runtime environment.

Runtime Environments

■ Need to consider

- What do objects look like in memory?
- What do functions look like in memory?
- Where in memory should they be placed?

■ There are no right answers to these questions.

- Many different options and tradeoffs.
- We will see several approaches.

Data Representations

- What do different types look like in memory?
- Machine typically supports only limited types:
 - Fixed-width integers: 8-bit, 16-bit- 32-bit, signed, unsigned, etc.
 - Floating point values: 32-bit, 64-bit, 80-bit IEEE 754.
- How do we encode our object types using these types?

Outline


- Introduction
- Runtime Environments
- **Encoding in Runtime Environments**
- Control-flow

Encoding Primitive Types

- Primitive integral types (byte, char, short, int, long, unsigned, uint16_t, etc.) typically map directly to the underlying machine type.
- Primitive real-valued types (float, double, long double) typically map directly to underlying machine type.
- Pointers typically implemented as integers holding memory addresses.
 - Size of integer depends on machine architecture; hence 32-bit compatibility mode on 64-bit machines.

Encoding Arrays

- C-style arrays:
 - Elements laid out consecutively in memory.



Encoding Arrays

■ Java-style arrays:


- Elements laid out consecutively in memory with size information prepended.



Encoding Arrays

■ D-style arrays:

- Elements laid out consecutively in memory
- Array variables store pointers to first and past-the-end elements.




Encoding Multidimensional Arrays

- Often represented as an array of arrays.
- Shape depends on the array type used.

Encoding Multidimensional Arrays

- Often represented as an array of arrays.
- Shape depends on the array type used.
- C-style arrays:


```
int a[3][2];
```



Encoding Multidimensional Arrays

- Often represented as an array of arrays.
- Shape depends on the array type used.
- Java-style arrays:

```
int[][] a = new int [3][2];
```



Encoding Functions

- Many questions to answer:
 - What does the dynamic execution of functions look like?
 - Where is the executable code for functions located?
 - How are parameters passed in and out of functions?
 - Where are local variables stored?

- The answers strongly depend on what the language supports.

The Procedure

- Procedures are the key to building large systems; they provide:
 - Control abstraction: well-defined entries & exits.
 - Name Space: has its own protected name space.
 - External Interface: access is by name & parameters.
- Requires system wide-compact:
 - Broad agreement on memory layout, protection, etc...
 - Must involve compiler, architecture, OS

The Procedure

- Establishes the need for private context:
 - Create a run-time “record” for each procedure to encapsulate information about control & data abstractions.
- Separate compilation:
 - Allows us to build large systems; keeps compile-time reasonable

The Procedure: a more detailed view

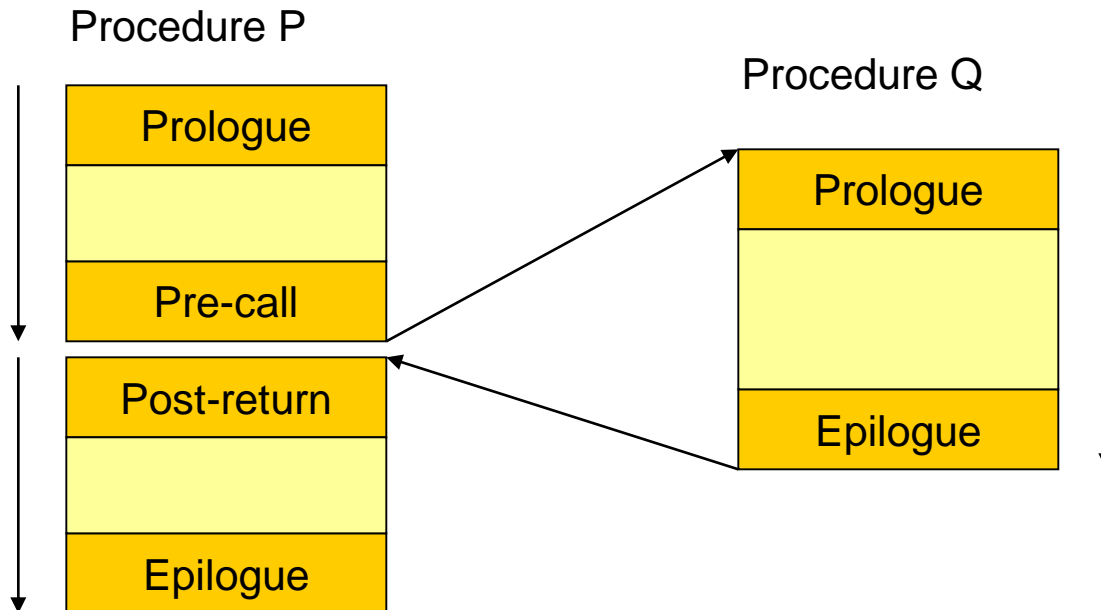
- A procedure is a collection of commands.
- The commands in a procedure are running line by line.

Outline

- Introduction
- Runtime Environments
- Encoding in Runtime Environments
- **Control-flow**

The linkage convention

- Procedures have well-defined control-flow behaviour:
 - A protocol for passing values and program control at procedure call and return is needed.
 - The linkage convention ensures that procedures inherit a valid run-time environment and that they restore one for their parents.
- Linkages execute at run-time.
- Code to make the linkage is generated at compile-time.

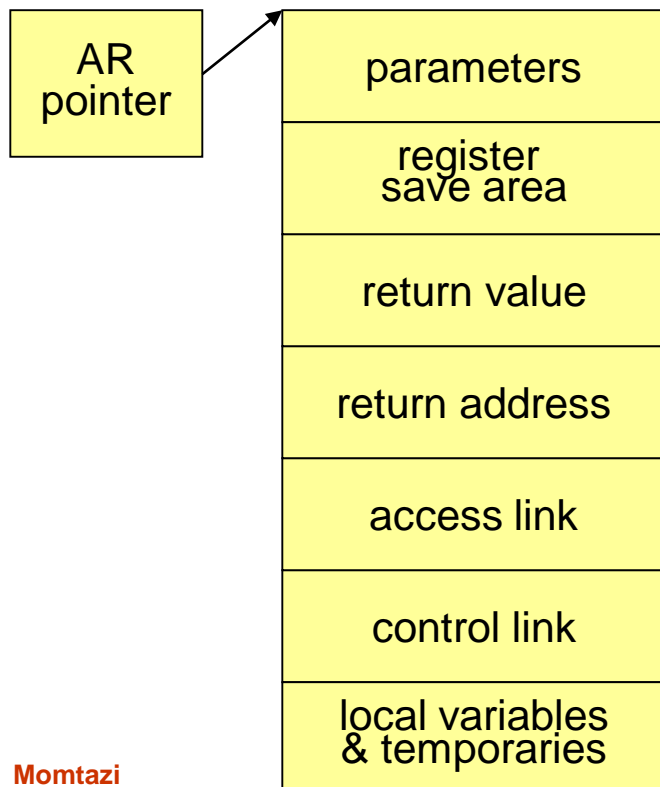


Control-flow as a Tree

- The control-flow in a program can be considered as an tree: Activation Tree.
- The root is the main program.
- The control-flow of the program is derived by search over this tree.

Storage Organisation: Activation Records

- Local variables require storage during the lifetime of the procedure invocation at run-time.
- The compiler arranges to set aside a region of memory for each individual call to a procedure (run-time support): Activation Record



In general, the compiler is free to choose any convention for the AR. The manufacturer may want to specify a standard for the architecture.

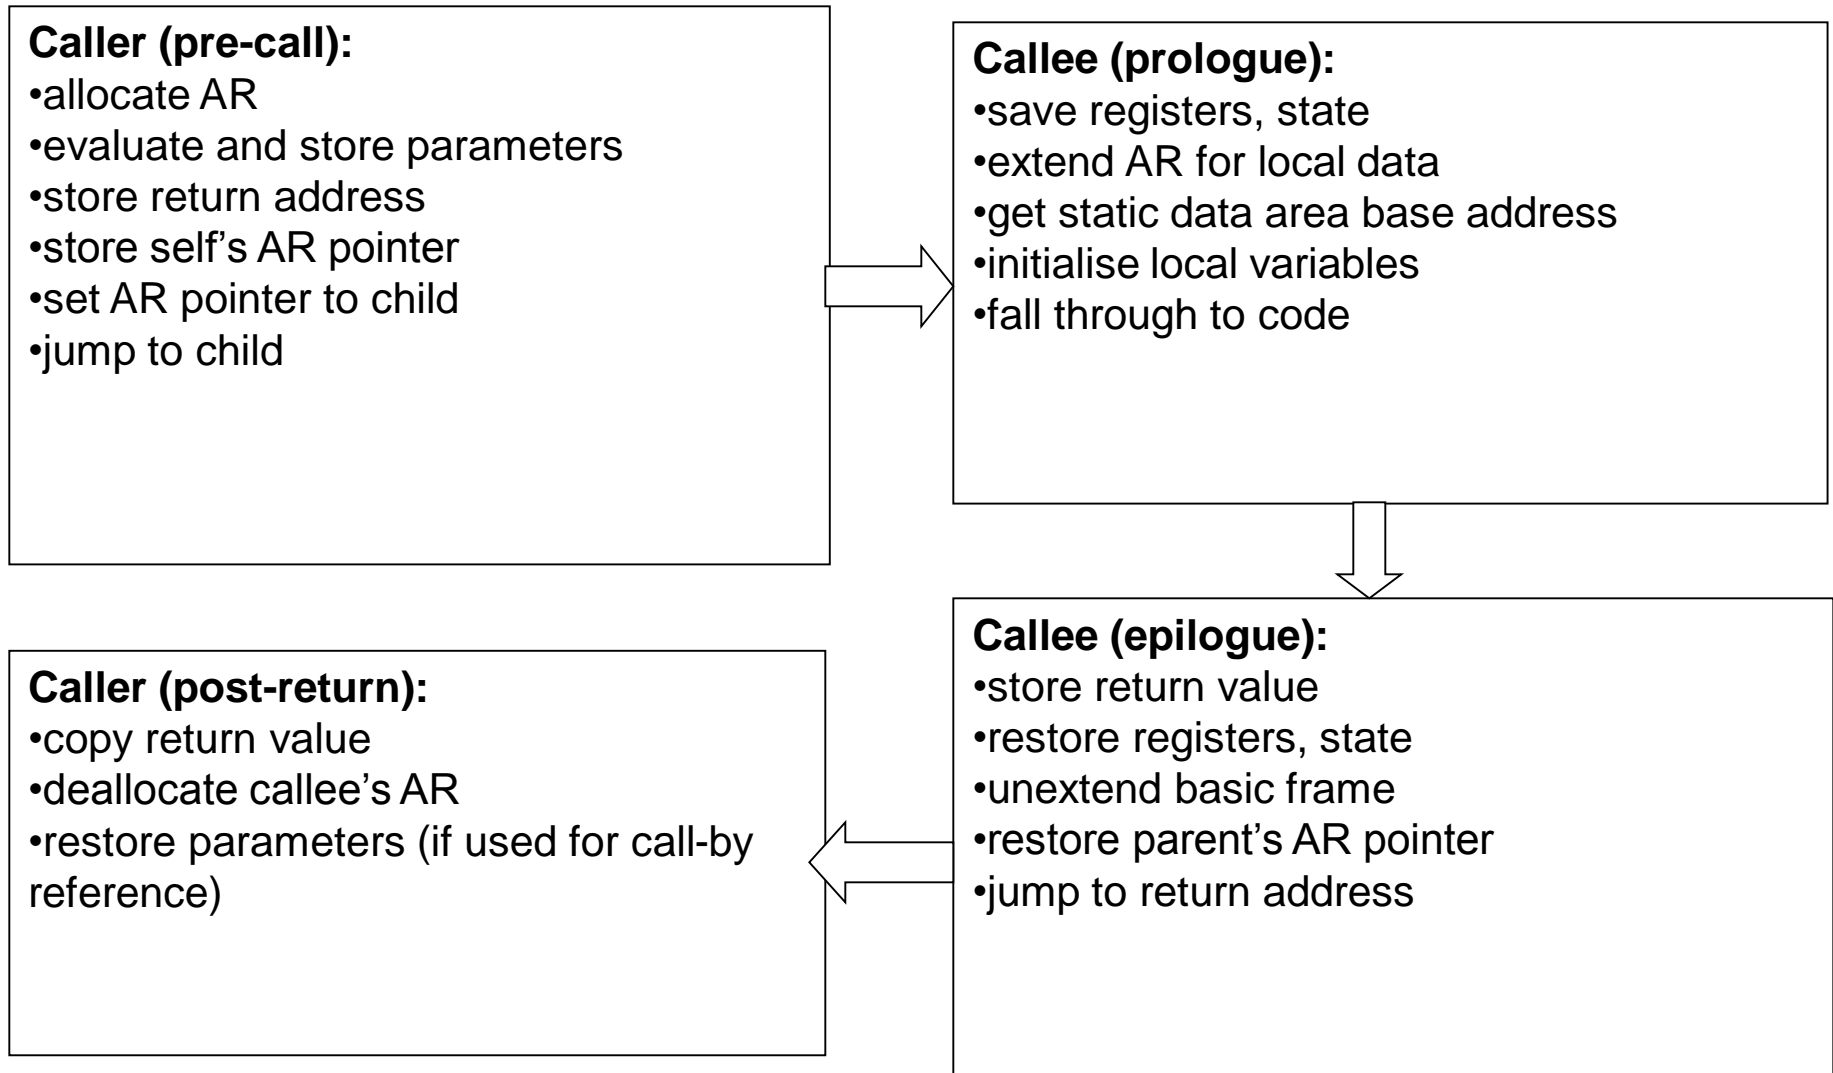
→ Address to resume caller

→ Help with non-local access

→ Pointer to caller's activation record

Activation Records are also known as stack frames.

Procedure linkages



the procedure linkage convention is a machine-dependent contract between the compiler, the OS and the target machines to divide clearly responsibility

Parameters in Procedure linkages

■ Parameters:

- Formal

- The arguments used in definition of a procedure

- Actual

- The arguments used when calling a procedure

■ When calling a procedure, the actual parameters are passed to the procedure. The formal parameters are replaced with the actual parameters.

```
int main() {  
    Fib(3);  
}  
  
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```

Activation Trees

- An activation tree is a tree structure representing all of the function calls made by a program on a particular execution.
 - Depends on the runtime behavior of a program; can't always be determined at compile-time.
 - (The static equivalent is the call graph).
- Each node in the tree is an activation record.
- Each activation record stores a control link to the activation record of the function that invoked it.

Activation Trees

```
int main() {  
    Fib(3);  
}  
  
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```

Activation Trees


```
int main() {  
    Fib(3);  
}  
  
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```



main

Activation Trees


```
int main() {  
    Fib(3);  
}  
  
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```



Activation Trees

```
int main() {  
    Fib(3);  
}
```


```
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```



Activation Trees

```
int main() {  
    Fib(3);  
}
```


```
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```



Activation Trees

```
int main() {  
    Fib(3);  
}
```


```
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```



Activation Trees


```
int main() {  
    Fib(3);  
}
```

```
int Fib(int n) {  
    if (n <= 1) return n;  
    return Fib(n - 1) + Fib(n - 2);  
}
```



Placing run-time data structures

Single logical address space:



- Code, static, and global data have known size.
 - They are defined at compile time.
- Heap & stack are used for allocating dynamic memory
 - They are defined at runtime
- Heap & stack grow towards each other
- The usage of heap and stack depends on the way the variables needs to be allocated: first-in-first-out or last-in-first-out

Stack of Activation Records

- Function calls are often implemented using a stack of activation records (or stack frames).
- Calling a function pushes a new activation record onto the stack.
- Returning from a function pops the current activation record from the stack.
- The runtime stack is an optimization of the activation tree.

Why Can We Optimize the Stack?

- Once a function returns, its activation record cannot be referenced again.
 - Every activation record has either finished executing or is an ancestor of the current activation record.
- We don't need to store old nodes in the activation tree.
 - We don't need to keep multiple branches alive at any one time.
- These are not always true!

Run-time storage organisation

- The compiler must ensure that each procedure generates an address for each variable that it references:
 - Static and Global variables:
 - Addresses are allocated statically (at compile-time). (relocatable)
 - Procedure local variables:
 - Put them in the activation record if: sizes are fixed and values are not preserved.
 - Dynamically allocated variables:
 - Usually allocated and deallocated explicitly.
 - Handled with pointers.

Links in Activation Records

- The access link of a function is to access variables from the outer scope.
 - Also called static link.
- The control link of a function is a pointer to the activation record of the function that called it.
 - Also called dynamic link.
 - Used to determine where to resume execution after the function returns.

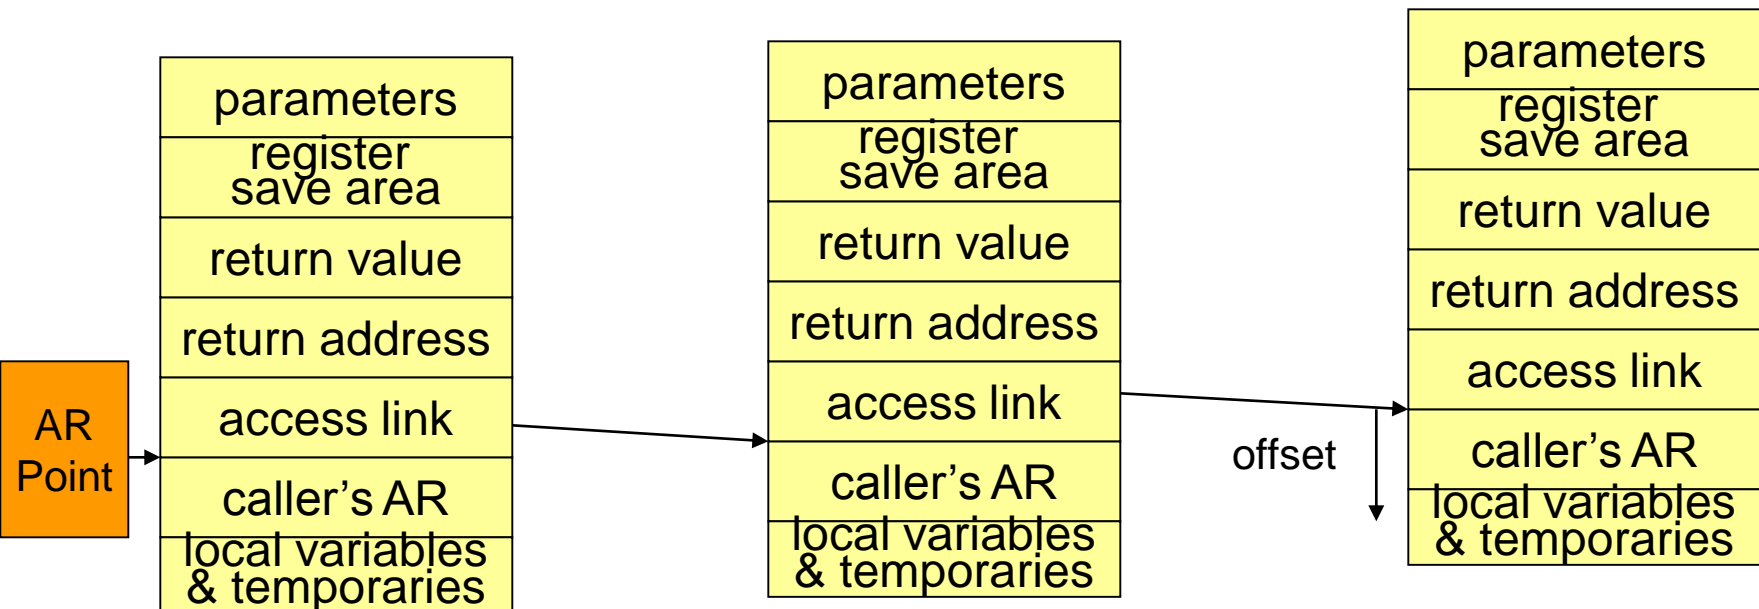
parameters
register save area
return value
return address
access link
control link
local variables & temporaries

Addressing non-local data

- In a language that supports nested lexical scopes, the compiler must provide a mechanism to map variables onto addresses.
- The compiler knows current level of lexical scope and of variable in question and offset (from the symbol table).
- Needs code to:
 - Track lexical ancestry (not necessarily the caller) among ARs.
 - Interpret difference between levels of lexical scope and offset.
- Two basic mechanisms:
 - Access links
 - Global display.

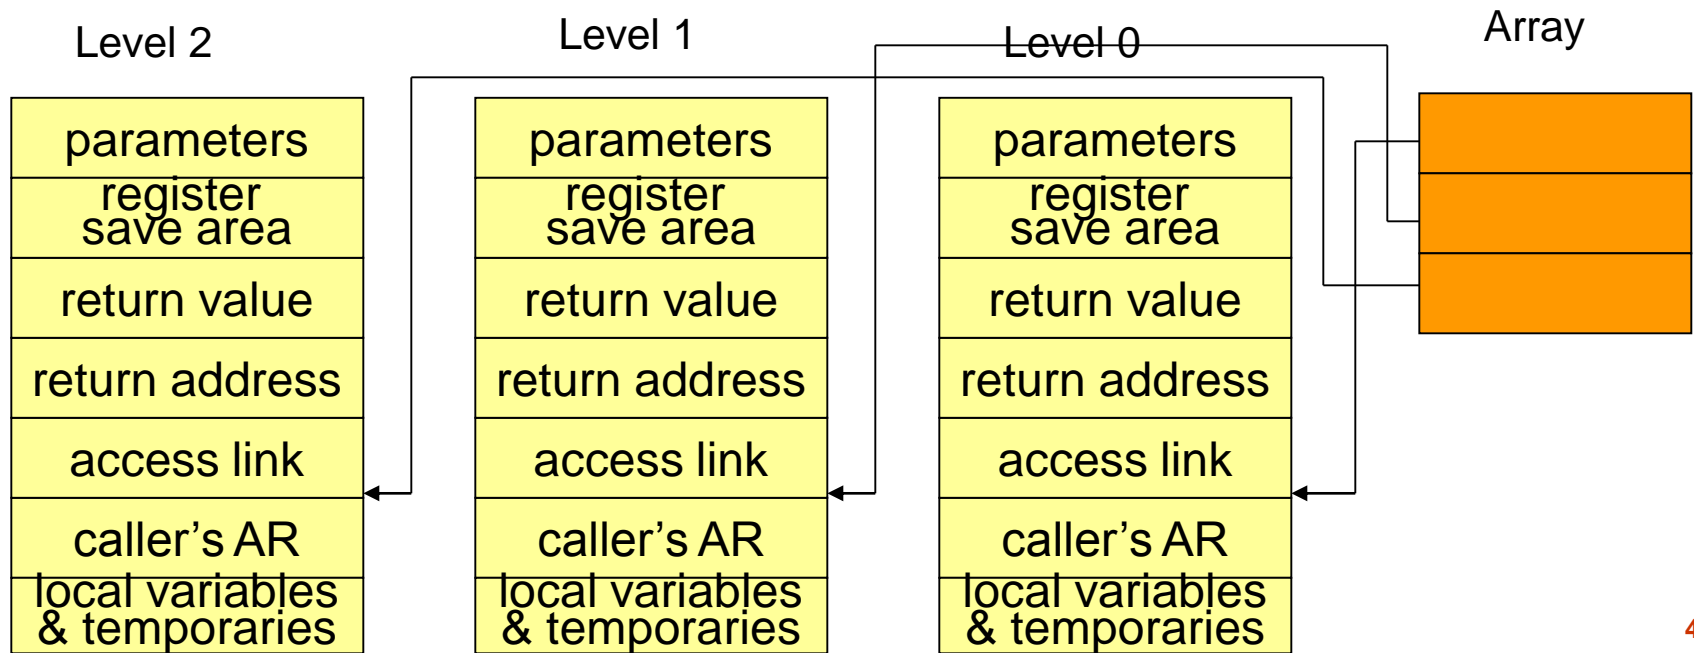
Access Links

- Idea: Each AR contains a pointer to its lexical ancestor.
- Compiler needs to emit code to find lexical ancestor (if caller's scope=callee's scope+1 then it is the caller; else walk through the caller's ancestors)
- Cost of access depends on depth of lexical nesting



Global Display

- Idea: keep a global array to hold ARPs for each level.
- Compiler needs to emit code (when calling and returning from a procedure) to maintain the array.
- Cost of access is fixed (table lookup + AR)
- Display vs access links trade-off. conventional wisdom: use access links when tight on registers; display when lots of registers.



Calling Sequence

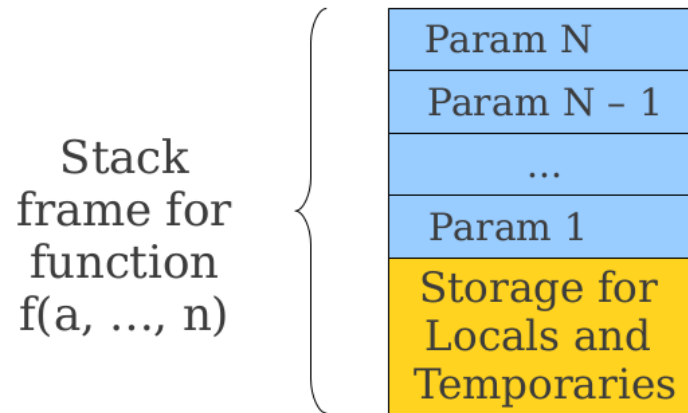
■ Assumptions

- “Once a function returns, its activation record cannot be referenced again.”
- “Every activation record has either finished executing or is an ancestor of the current activation record.”

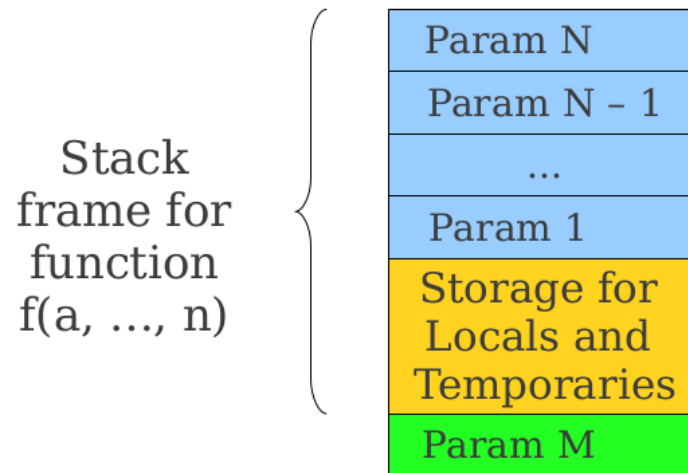
Calling Sequence

- Caller responsible for pushing and popping space for callee's arguments.
- Callee responsible for pushing and popping space for its own temporaries.
- Example:
 - Function `f()` is caller and function `g()` is callee.

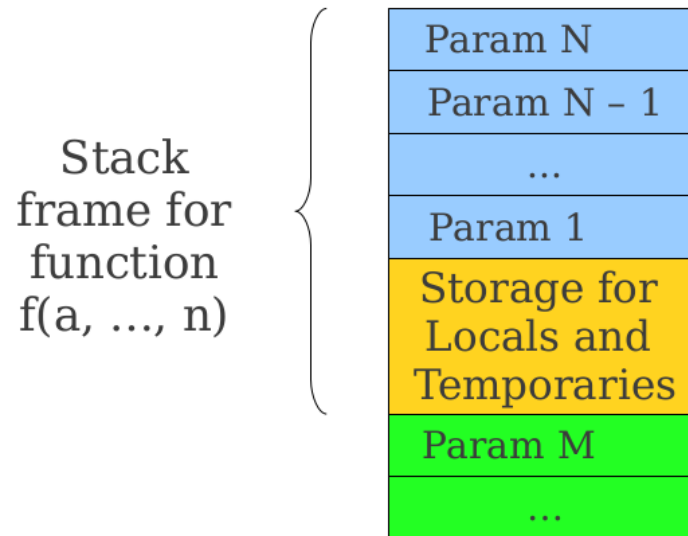
Calling Sequence



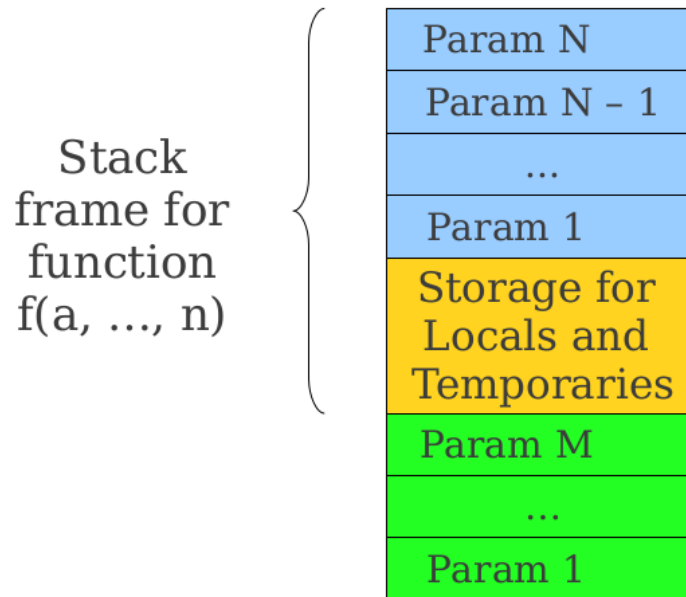
Calling Sequence



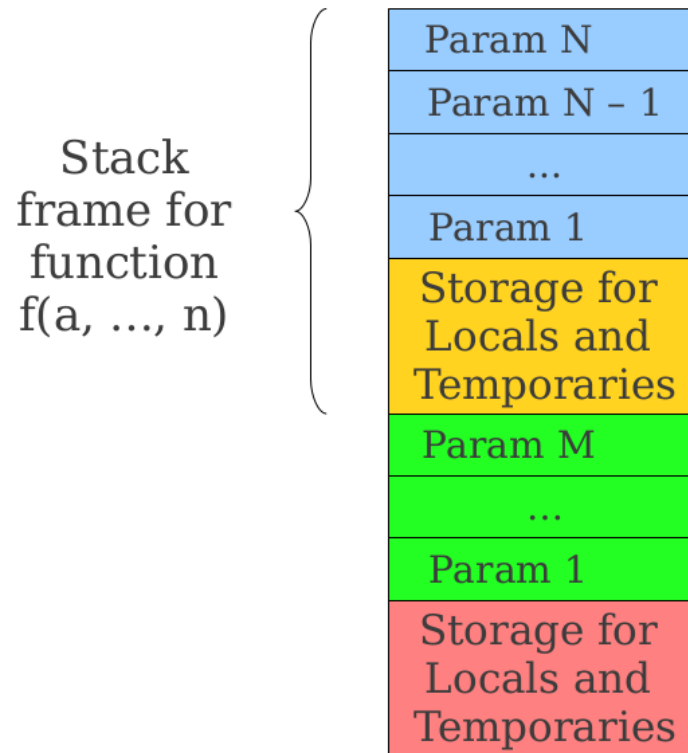
Calling Sequence



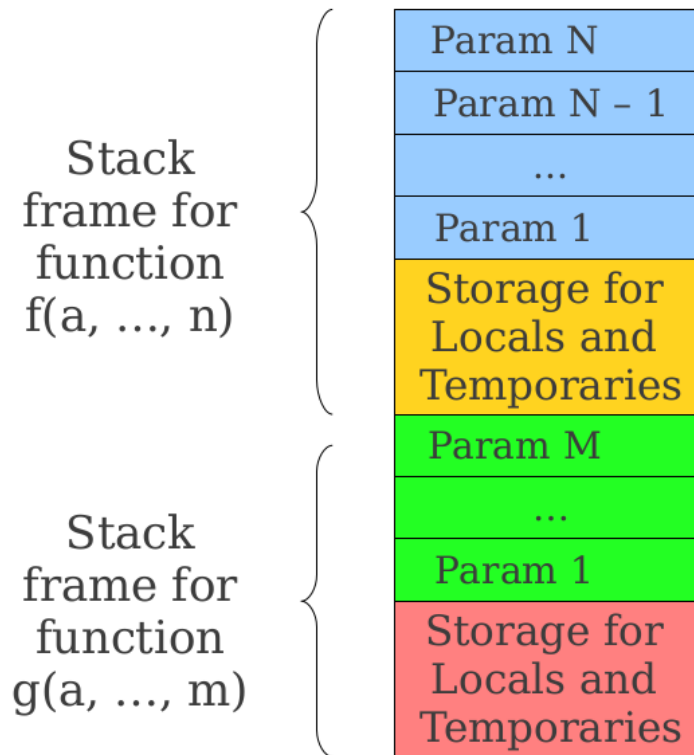
Calling Sequence



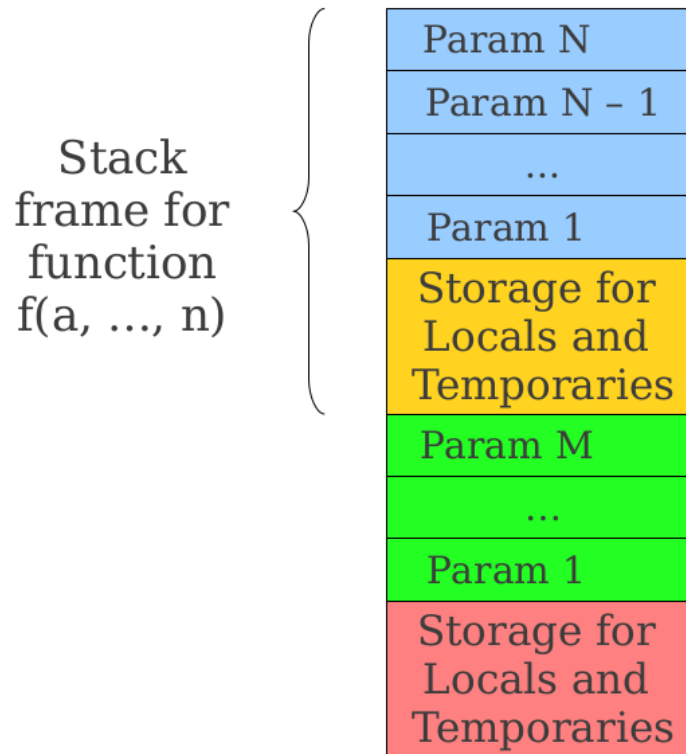
Calling Sequence



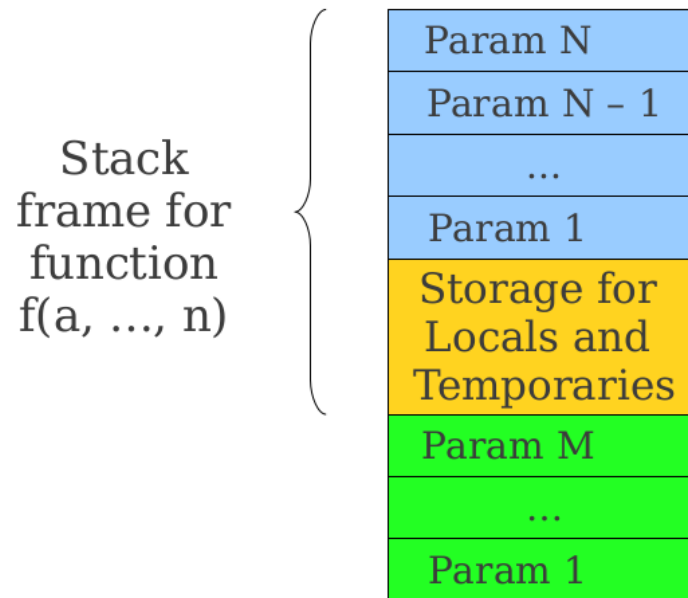
Calling Sequence



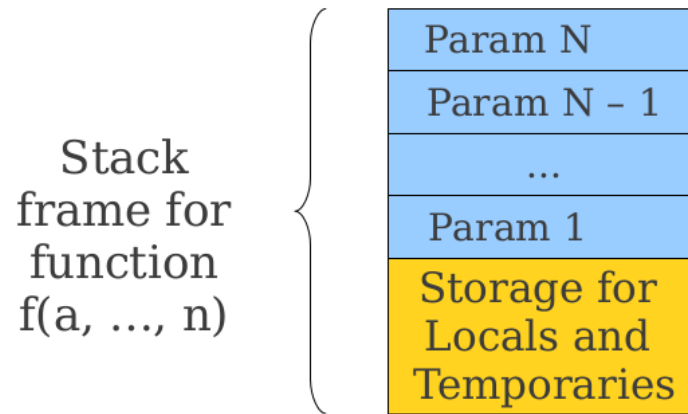
Calling Sequence



Calling Sequence



Calling Sequence



Example

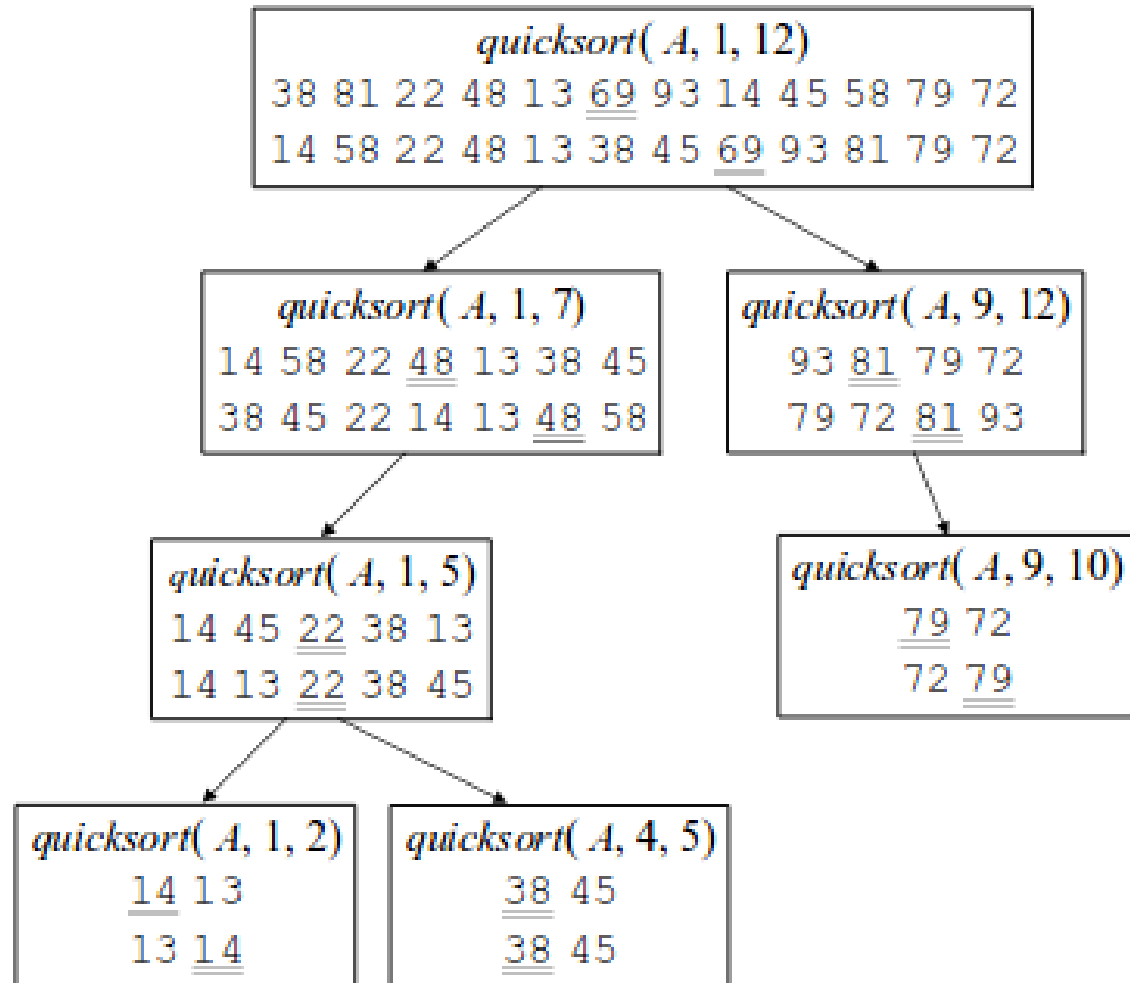
■ Quicksort Code:

```
/* low --> Starting index, high --> Ending index */
quicksort(arr[], low, high)
{
    if (low < high)
    {
        /* pi is partitioning index, arr[p] is now
           at right place */
        pi = partition(arr, low, high);

        quicksort(arr, low, pi - 1); // Before pi
        quicksort(arr, pi + 1, high); // After pi
    }
}
```

Example

■ Data:



Example

■ Quicksort Execution:

- push quicksort (1,12)
- push partition (1,12)
- pop partition (1,12)
- ...

Summary

- An activation records arranges to set aside a region of memory for each individual call to a procedure
- An Activation tree is a tree structure representing all of the function calls made by a program on a particular execution.
- The runtime stack is an optimization of the activation tree stack.
- Activation records logically store a control link to the calling function and an access link to the function in which it was created.

Question?